



Building a Single Logical Data Lake

*For Advanced Analytics, Data Science, and
Business Operations*

Ravi Shankar, CMO
September, 2018

denodo 

- Competition from a low cost vendor
- Lower the price, affecting margins?
- Or, maintain high price, but differentiate in other ways?




Large Heavy Equipment Manufacturer

Self-service / Predictive Analytics – IoT Integration



Tableau: Dealer / Customer Dashboard

denodo  DATA VIRTUALIZATION



Benefits

- Improved asset performance and proactive maintenance
- Increased revenue from sale of services and parts
- Reduced warranty costs of parts failure

“Big Data Challenges Impede Big Data Vision”



Base: 3,343 global data and analytics technology decision makers

Note: Multiple responses accepted; top eight responses shown.

Source: Forrester Data Global Business Technographics® Data And Analytics Survey, 2016 4

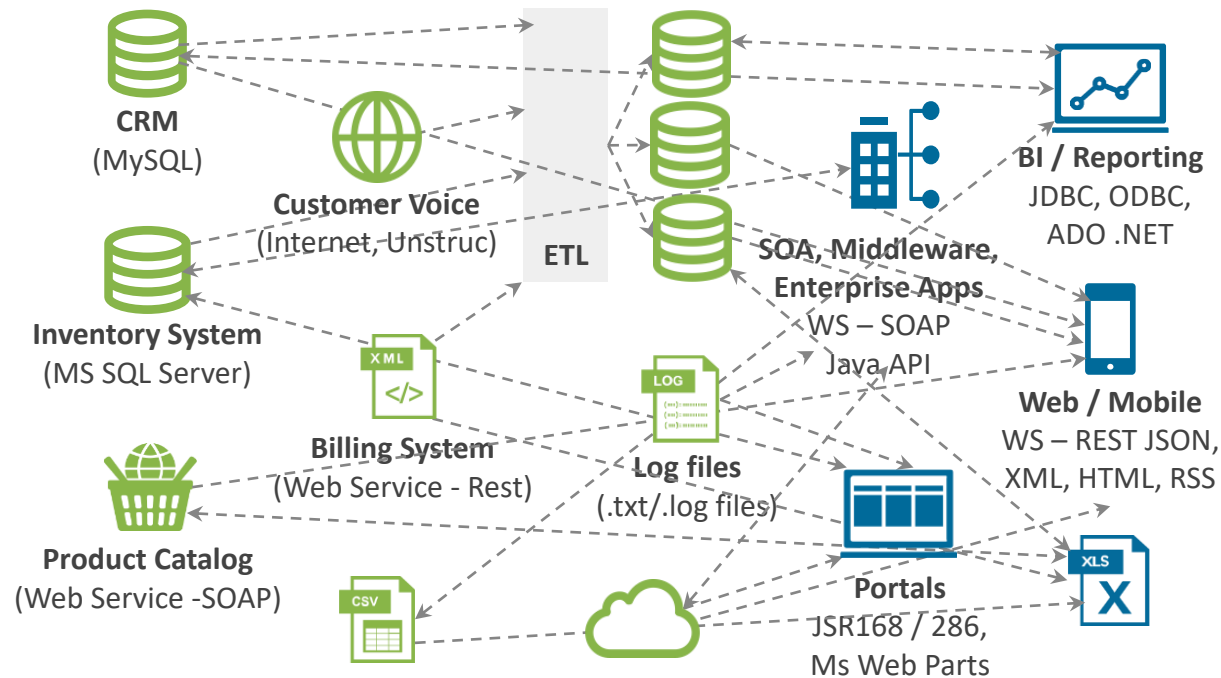
IT – Business Dilemma

IT Architecture is Unmanageable & Brittle because:

IT Focuses on
Data Collection
& Storage

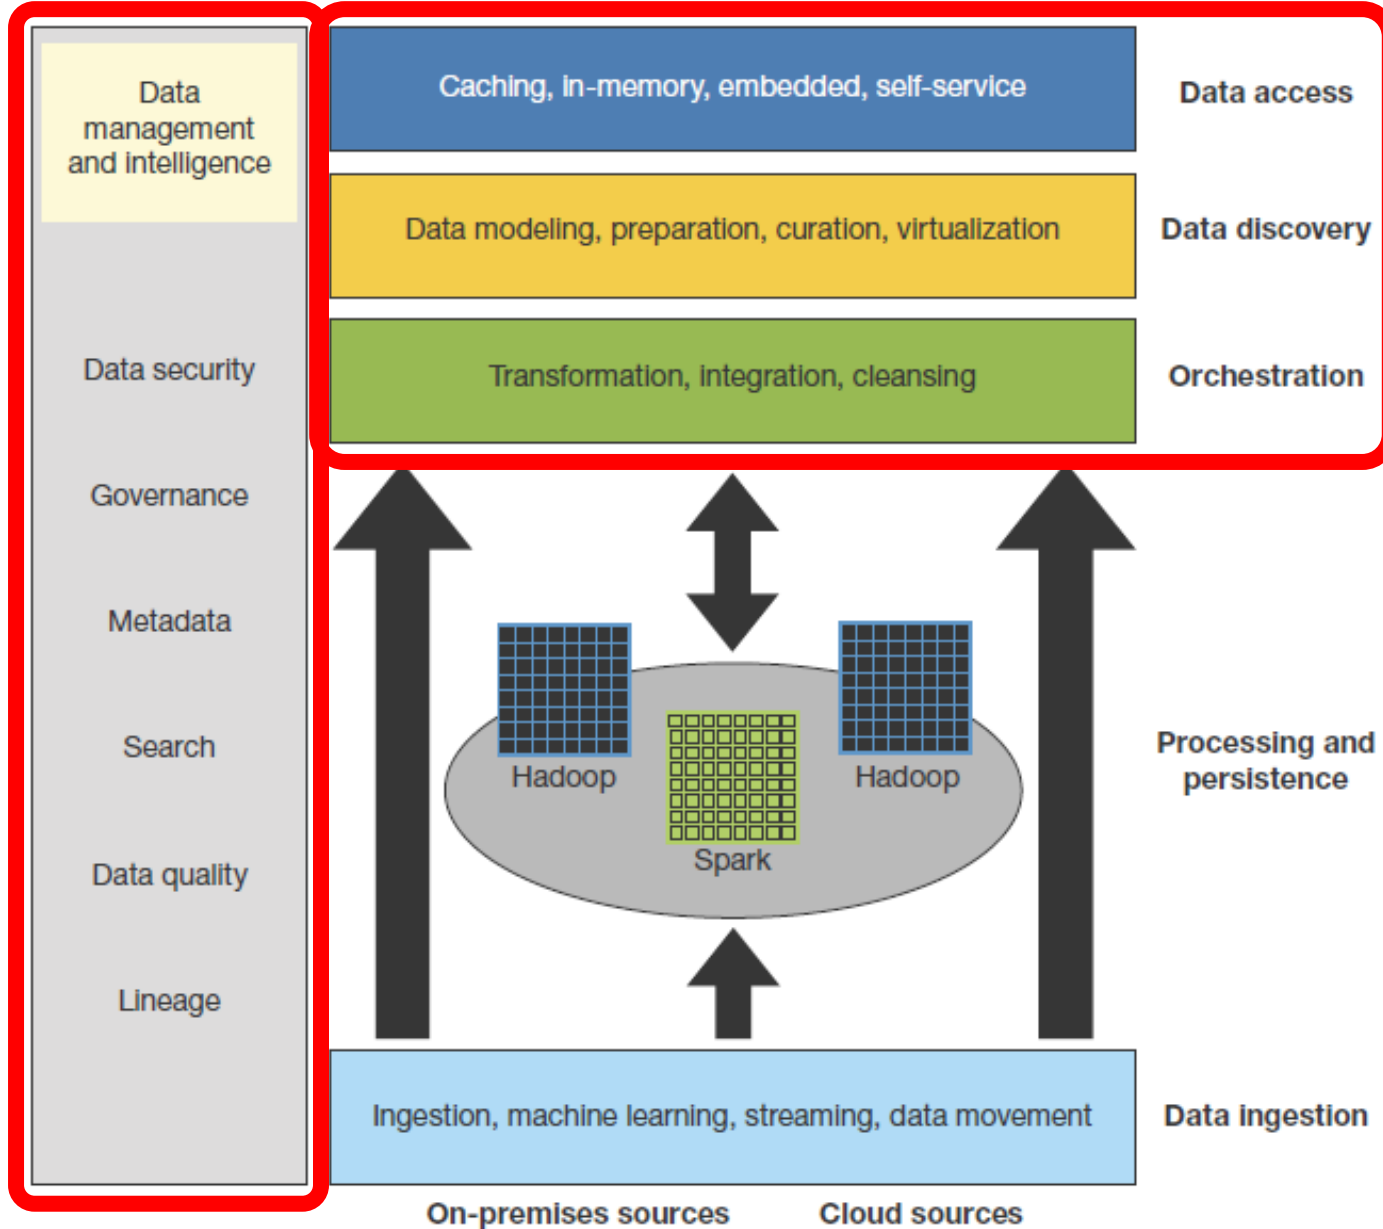
No One Focused on Data Delivery

– So create 100's to 1K's of brittle direct connections and replicate large volumes of data

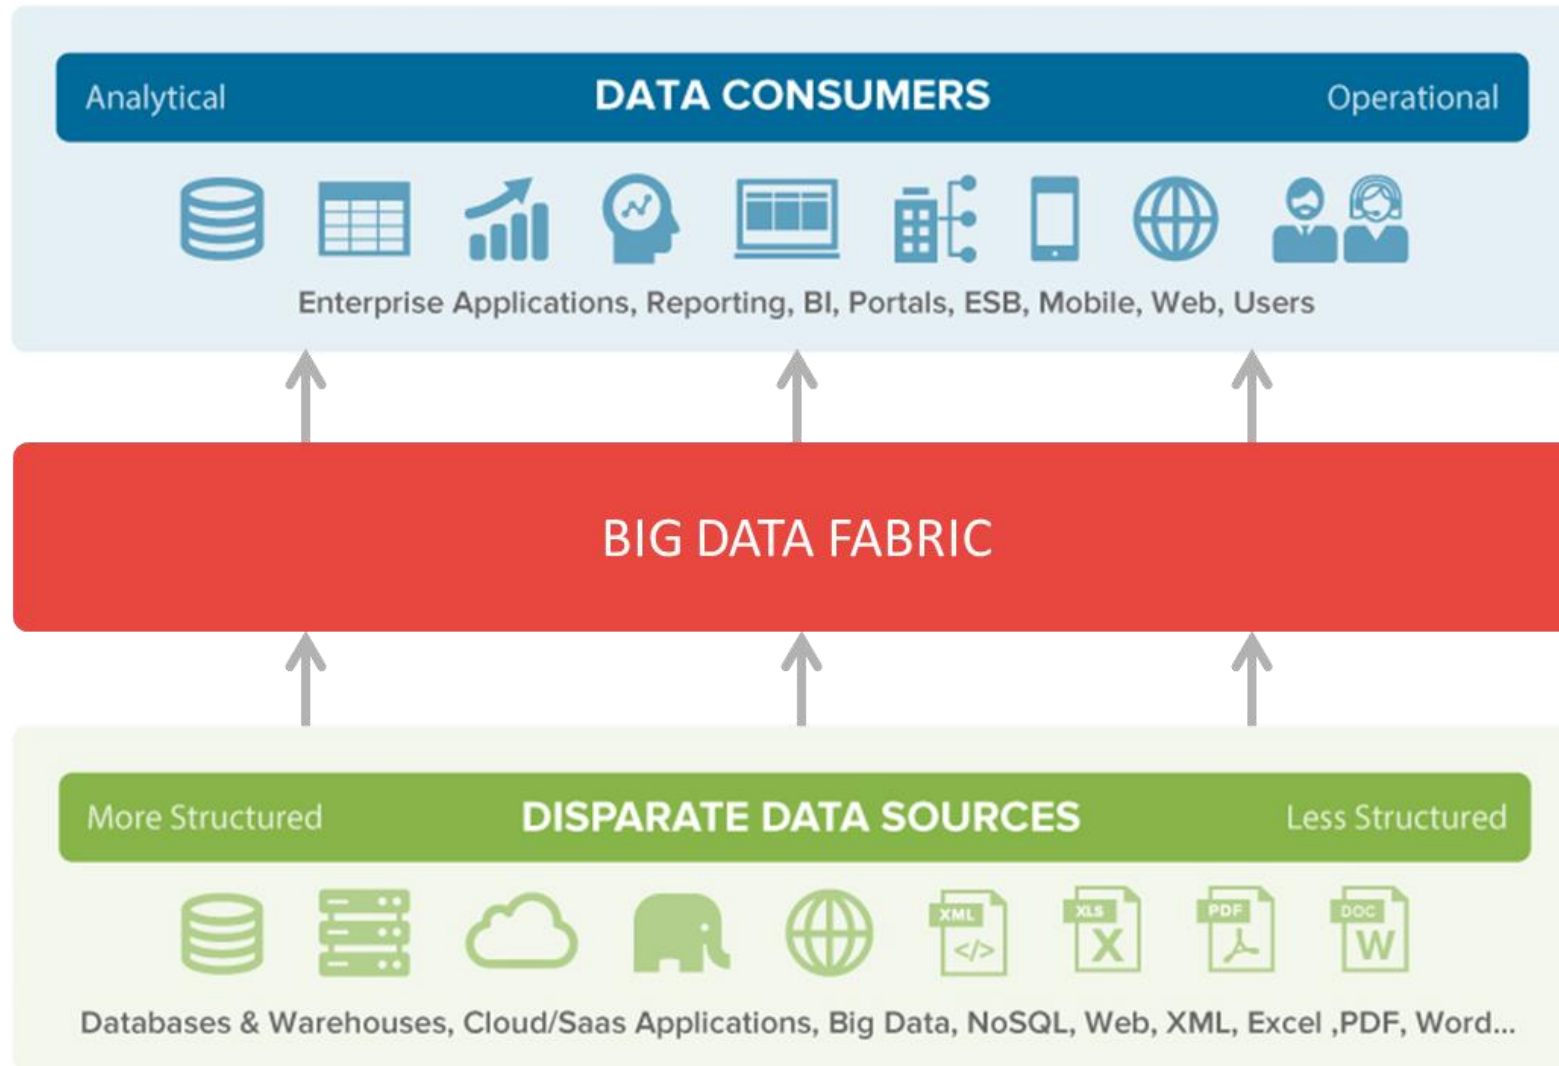


Business
Focuses on Data
Visualization &
Analysis

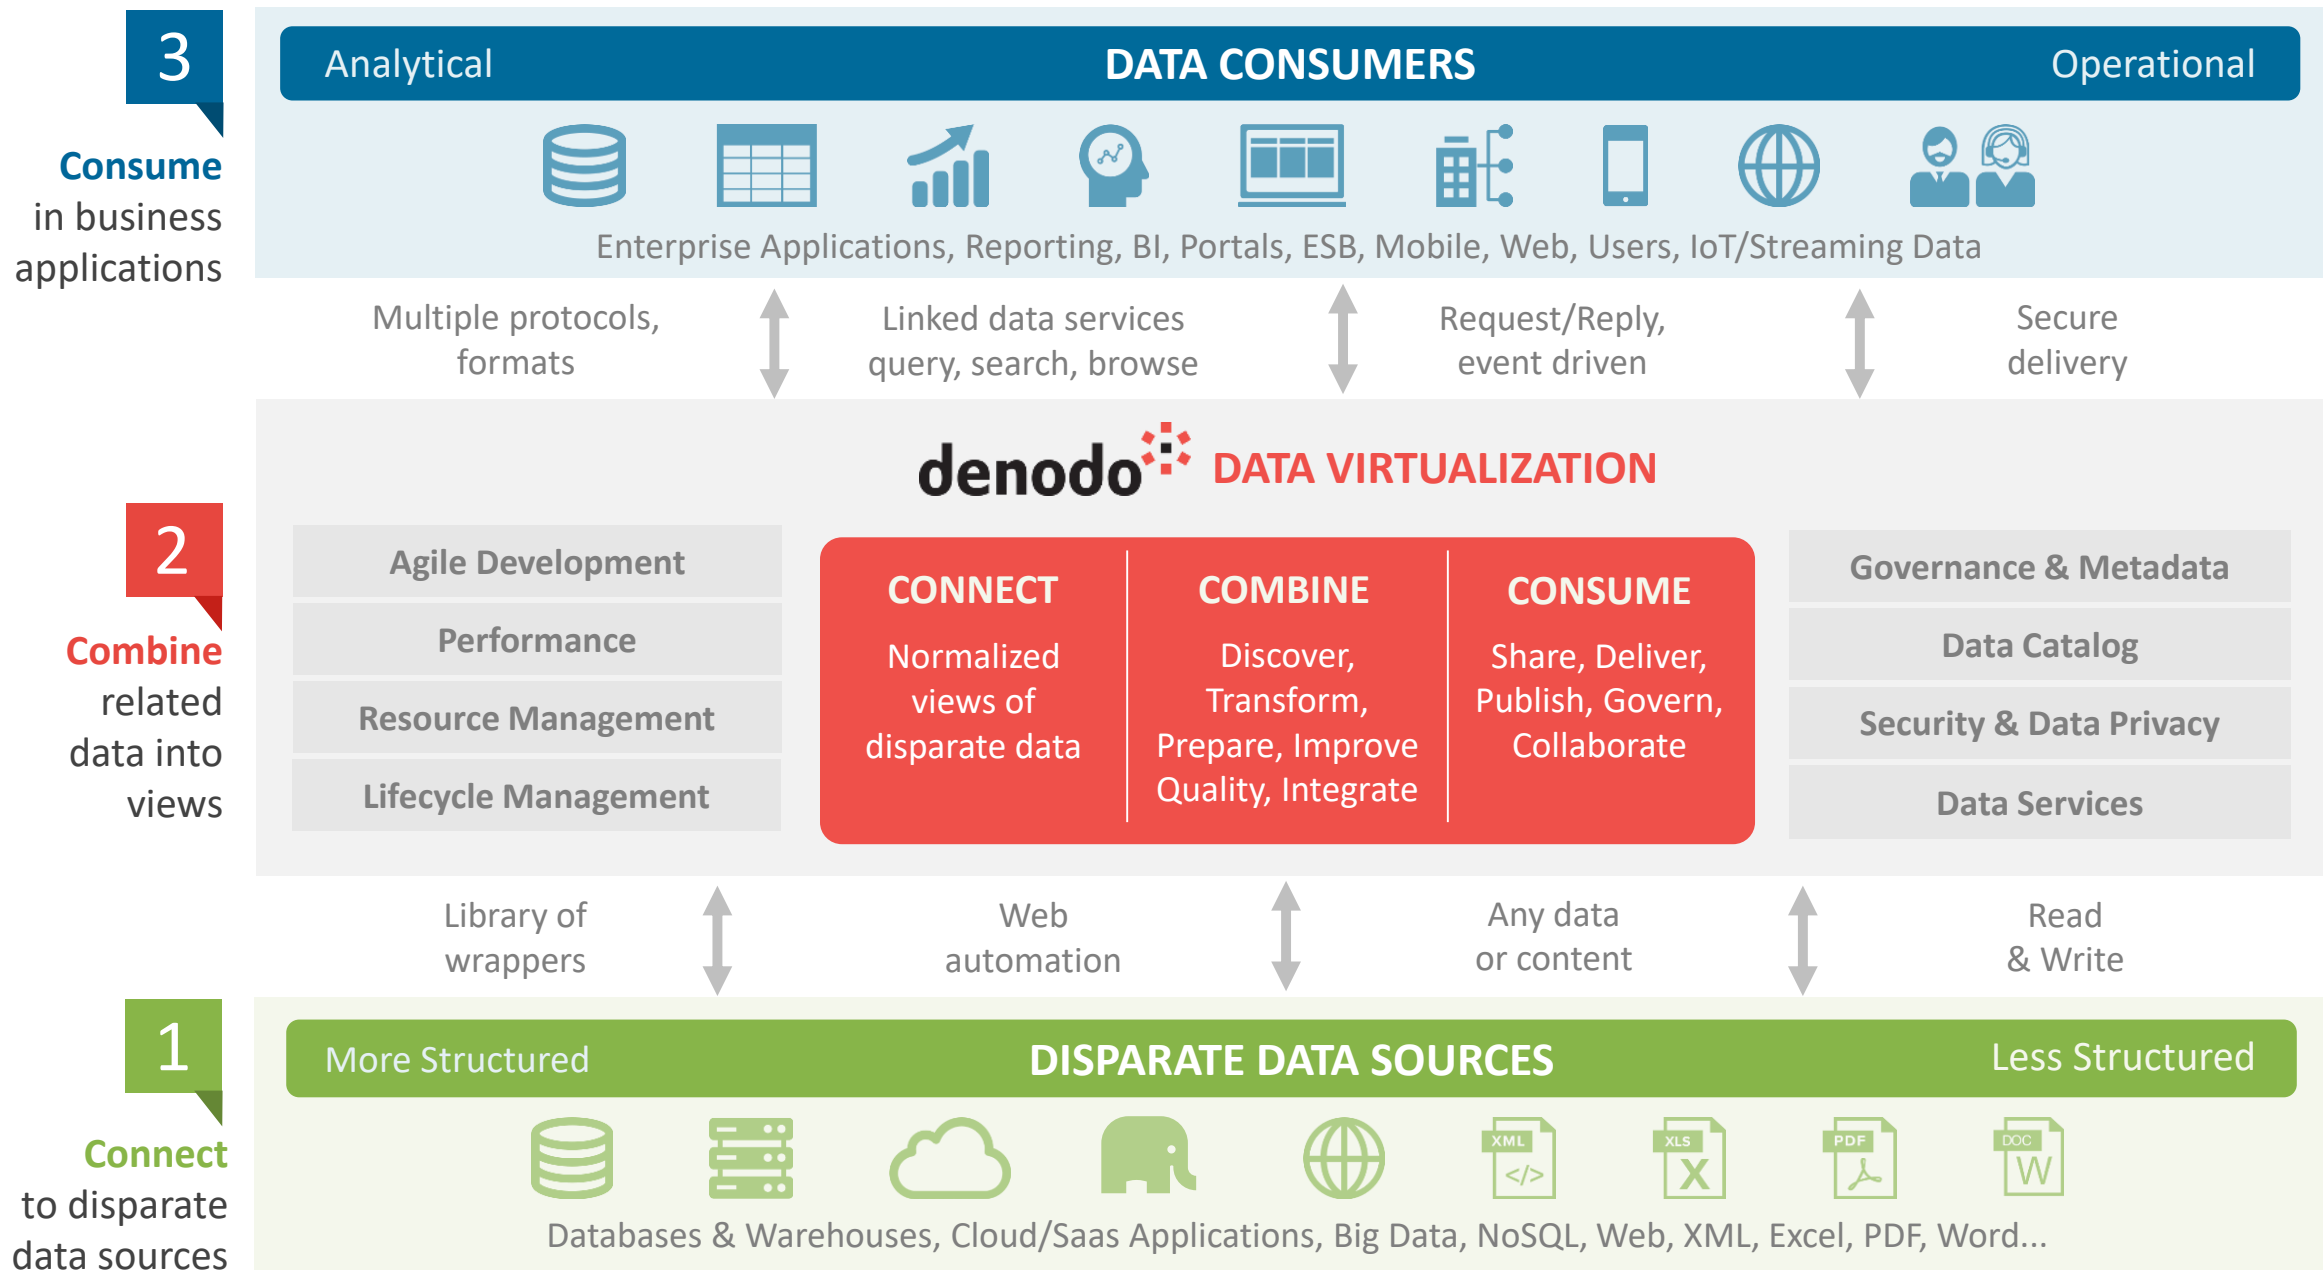
Big Data Fabric Architecture – Forrester Research, 2016



Big Data Fabric – Data Abstraction Layer

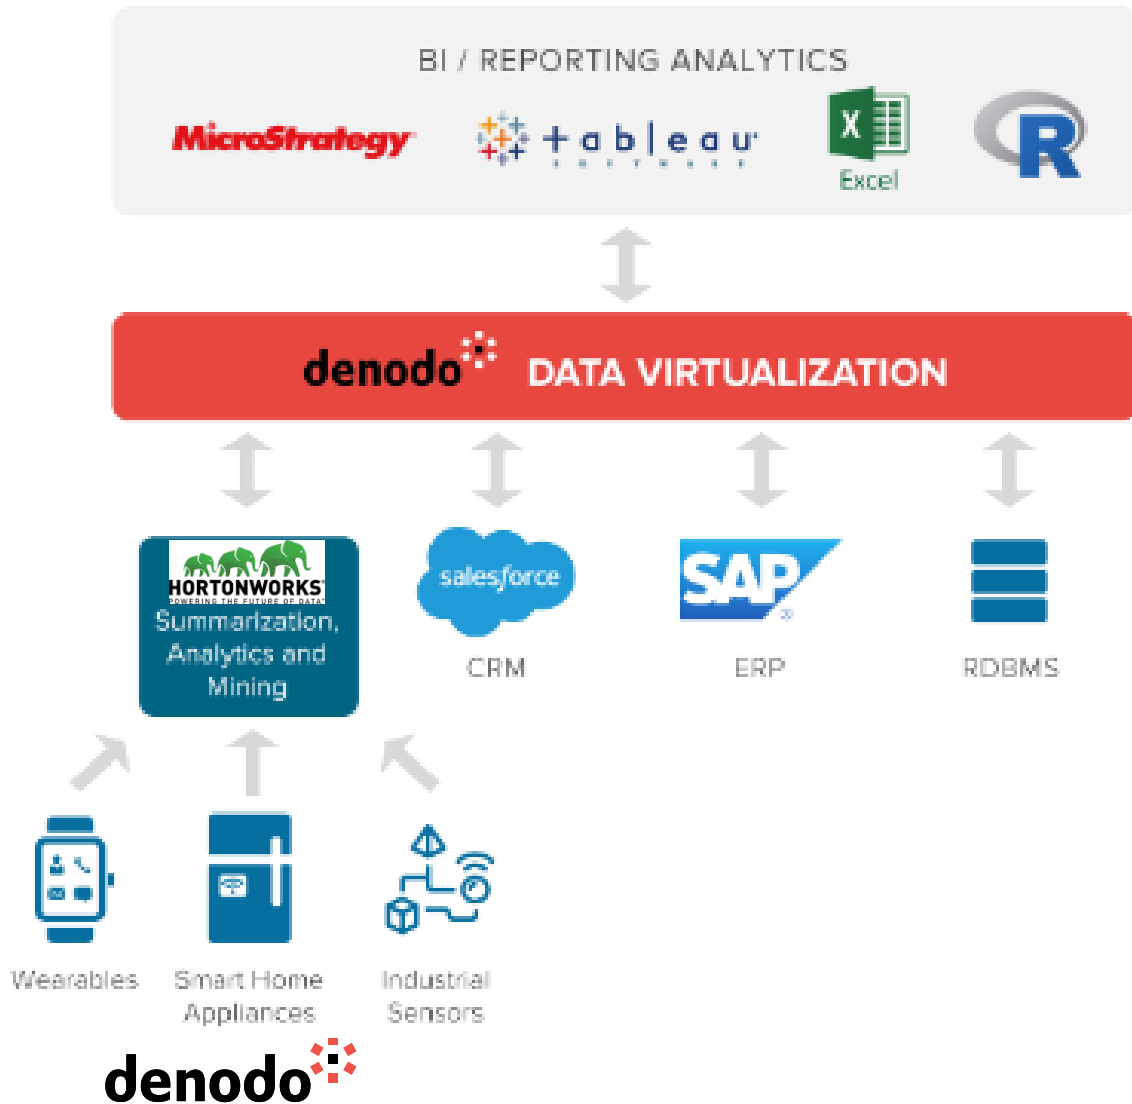


- | Abstracts access to **disparate data sources**
- | Acts as a **single repository** (virtual)
- | Makes data available in **real-time** to consumers

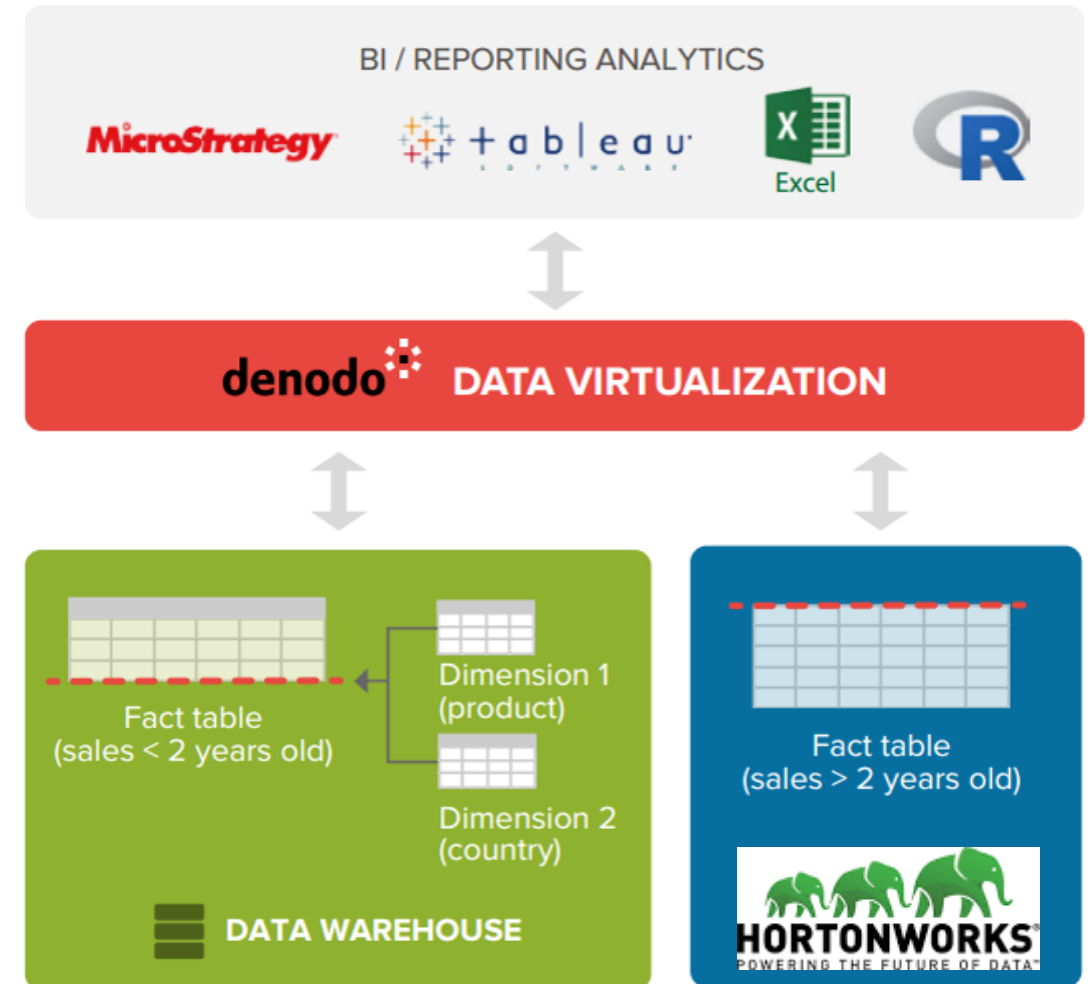


Logical Data Lake – Use Cases

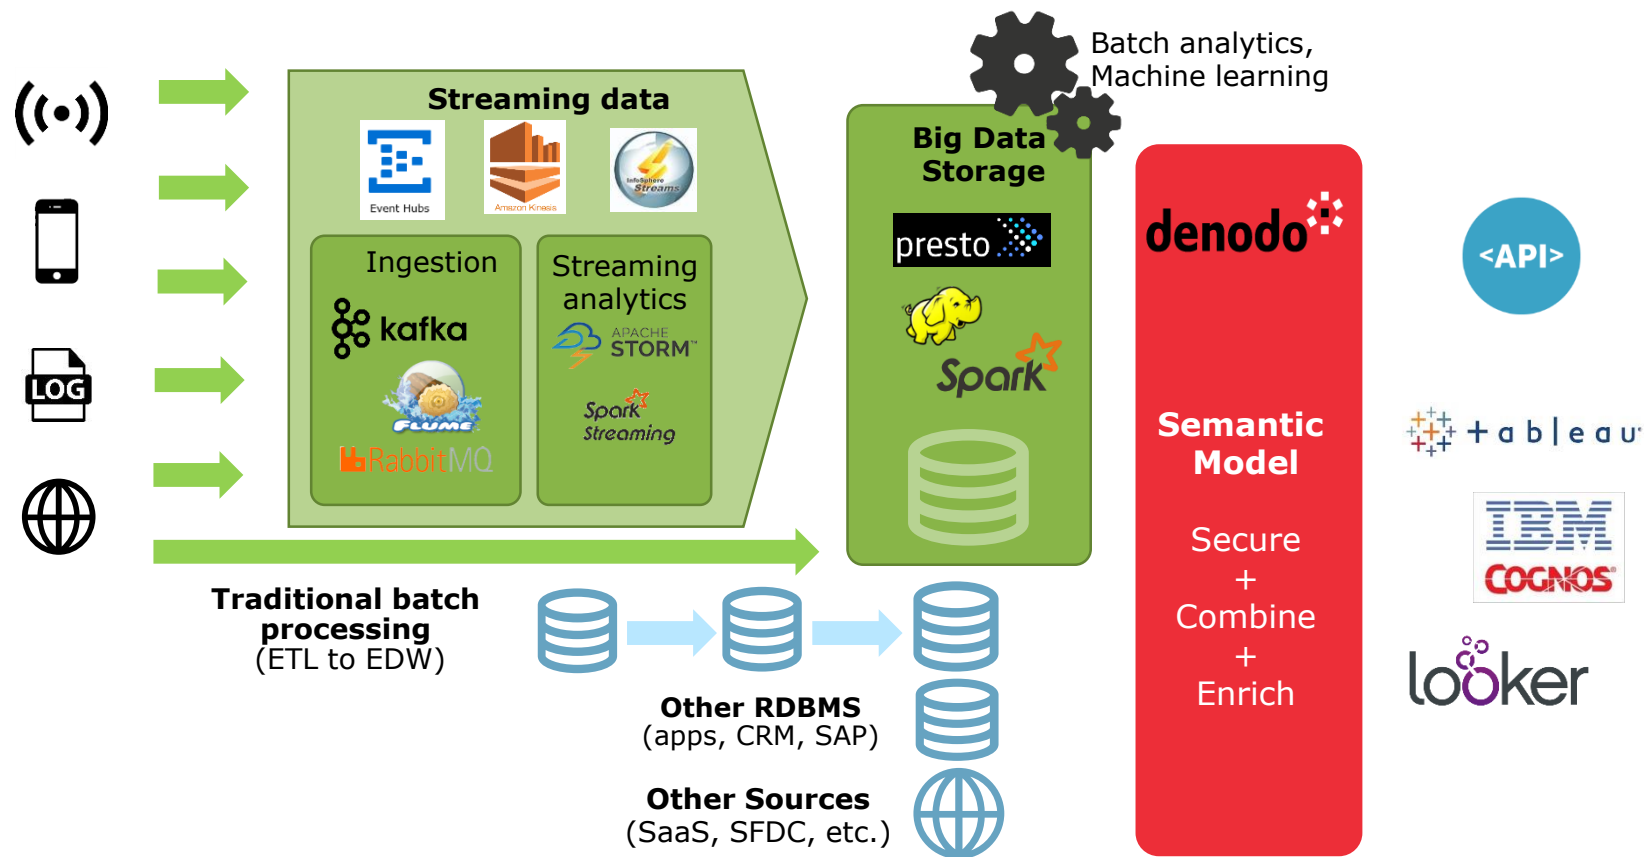
IoT Integration



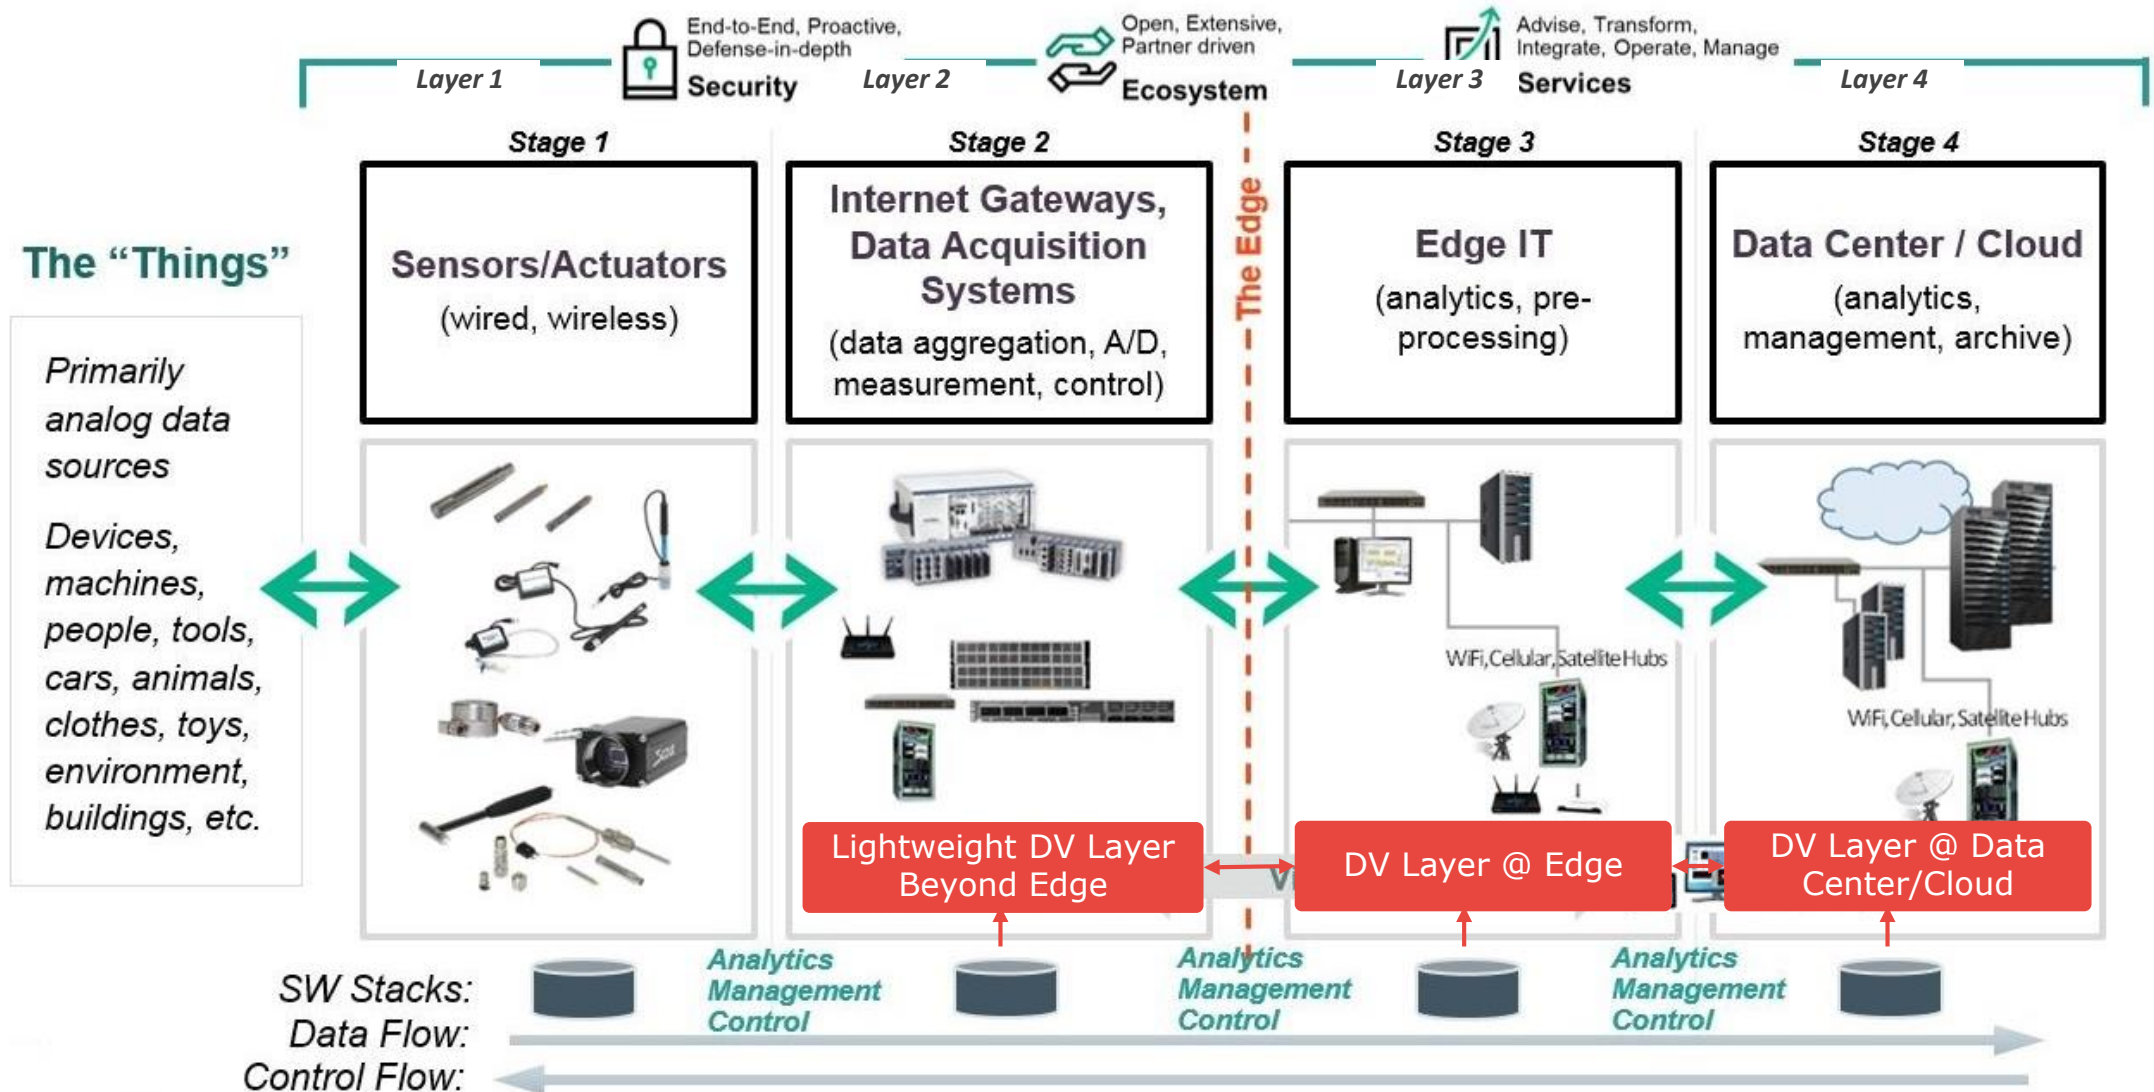
Data Warehouse Offloading



Data Virtualization in the IoT Ecosystem



Edge Computing Using Data Virtualization



Six Essential Capabilities of Data Virtualization

1. Data abstraction
2. Zero replication, zero relocation
3. Real-time information



4. Self-service data services
4. Centralized metadata, security & governance
4. Location-agnostic architecture for multi-cloud, hybrid acceleration

Big Data Queries Faster with Denodo Platform

Performance comparison of 5 different queries


1. Data Virtualization delivers better performance without needing to replicate data into Hadoop.
2. Data Virtualization leverages Data Source Architectures for what they are good at.

	Impala Hadoop-only Runtime (s)	Denodo Runtime (s)	Denodo Runtime w/ Cache (s)	Data Volumes
Query 1	199	120	68	Queries 1,2,3,5 <ul style="list-style-type: none">•Exadata Row Count: ~5M•Impala Row Count: ~500k Query 4 <ul style="list-style-type: none">•Exadata Row Count: ~5M•Impala Row Count: ~2M
Query 2	187	96	88	
Query 3	120	212	115	
Query 4	timeout	328	69	
Query 5	46	91	56	

Denodo Dynamic Query Optimizer

```
SELECT c.id, SUM(s.amount) as total
FROM customer c JOIN sales s
ON c.id = s.customer_id
GROUP BY c.id
```

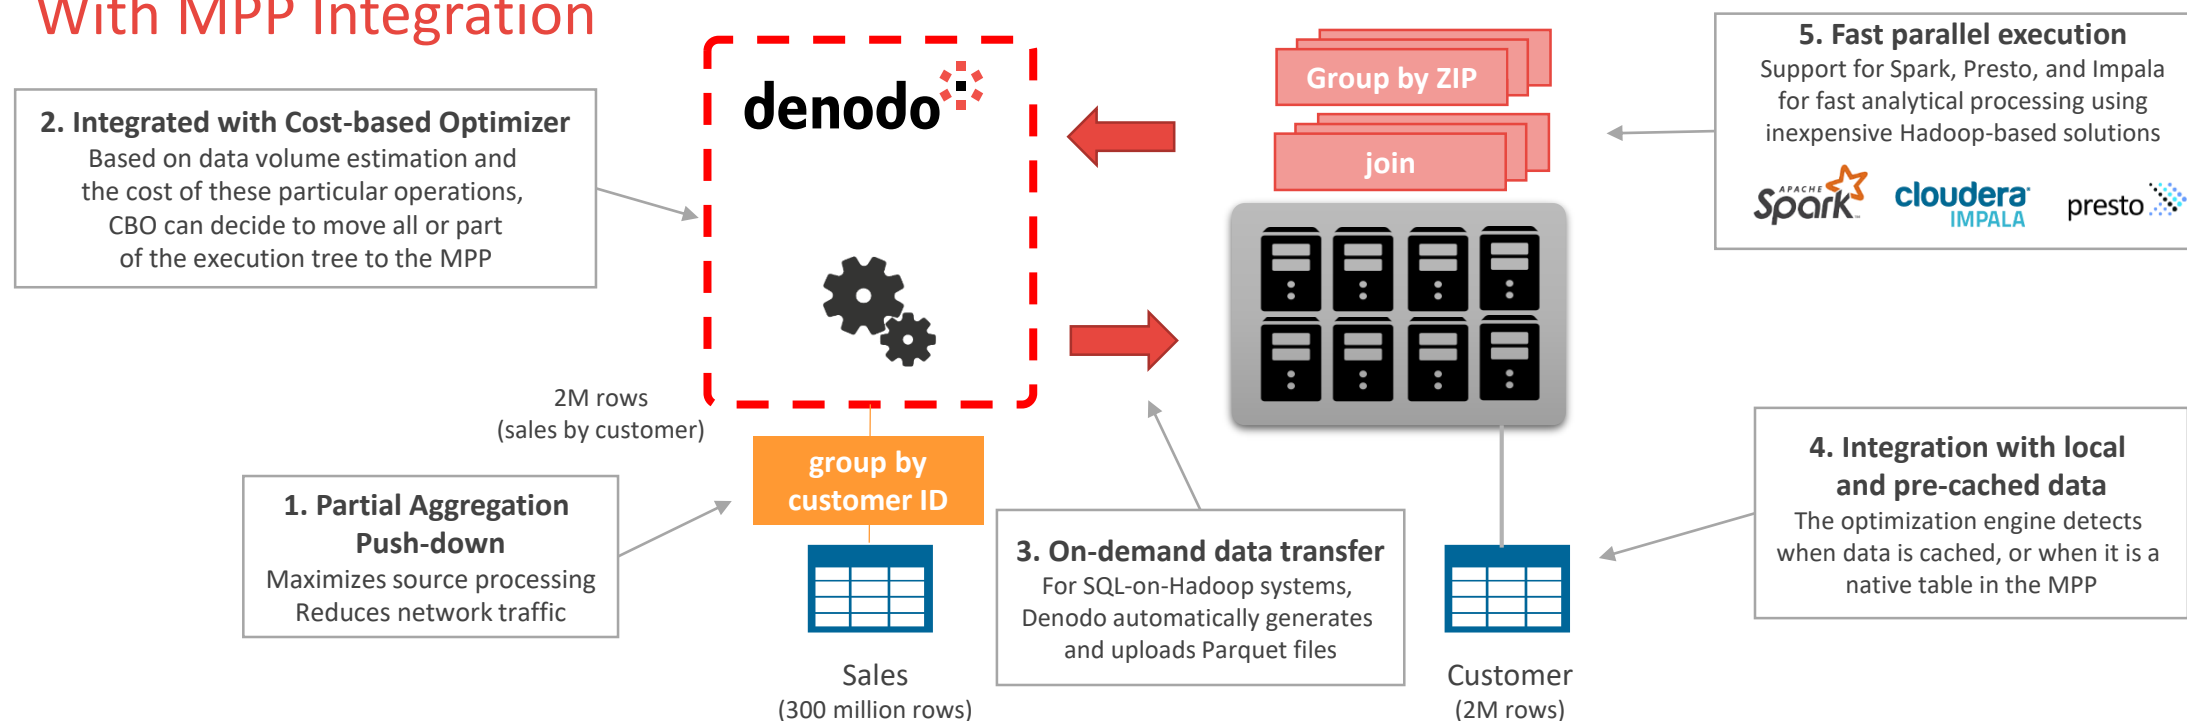
System	Execution Time	Data Transferred	Optimization Technique
Denodo	9 sec.	4 M	Aggregation push-down
Without Denodo	125 sec.	292 M	None: full scan



Data Virtualization with Massively Parallel Processing



With MPP Integration



System	Execution Time	Optimization Techniques
Others	~ 10 min	Traditional Federation Tools
No MPP	43 sec	Aggregation push-down
With MPP	11 sec	Aggregation push-down + MPP integration (Impala 8 nodes)

ROI and TCO of Data Virtualization

Customer-reported projected savings by percentage

Data Integration Cost reduction

- **60-80%** savings

Traditional Call Centres, Portals

- **30-70%** savings

BI and Reporting

- **40-60%** savings

ETL and Data Warehousing

- Project timelines of 6-12 months reduced to **3-6 months**
- **Up to 85%** reduction in time

- New sources can be configured in minutes, and fully integrated within days.
- **100's** of application entities can be integrated within weeks.
- New business functionality can be added within days.
- Existing functionality can be enhanced with new data within days.
- Data proliferation can be significantly reduced.
- Common, consistent and timely access to all data via preferred visualization tools.



Denodo 'Horizontal Solution' Categories



Customer Centricity / MDM

- ✓ Complete View of Customer



Data Governance

- ✓ GRC
- ✓ GDPR
- ✓ Data Privacy / Masking



Data Services

- ✓ Data as a Service
- ✓ Data Marketplace
- ✓ Data Services
- ✓ Application and Data Migration



BI and Analytics

- ✓ Self-Service Analytics
- ✓ Logical Data Warehouse
- ✓ Enterprise Data Fabric



Big Data

- ✓ Logical Data Lake
- ✓ Data Warehouse Offloading
- ✓ IoT Analytics



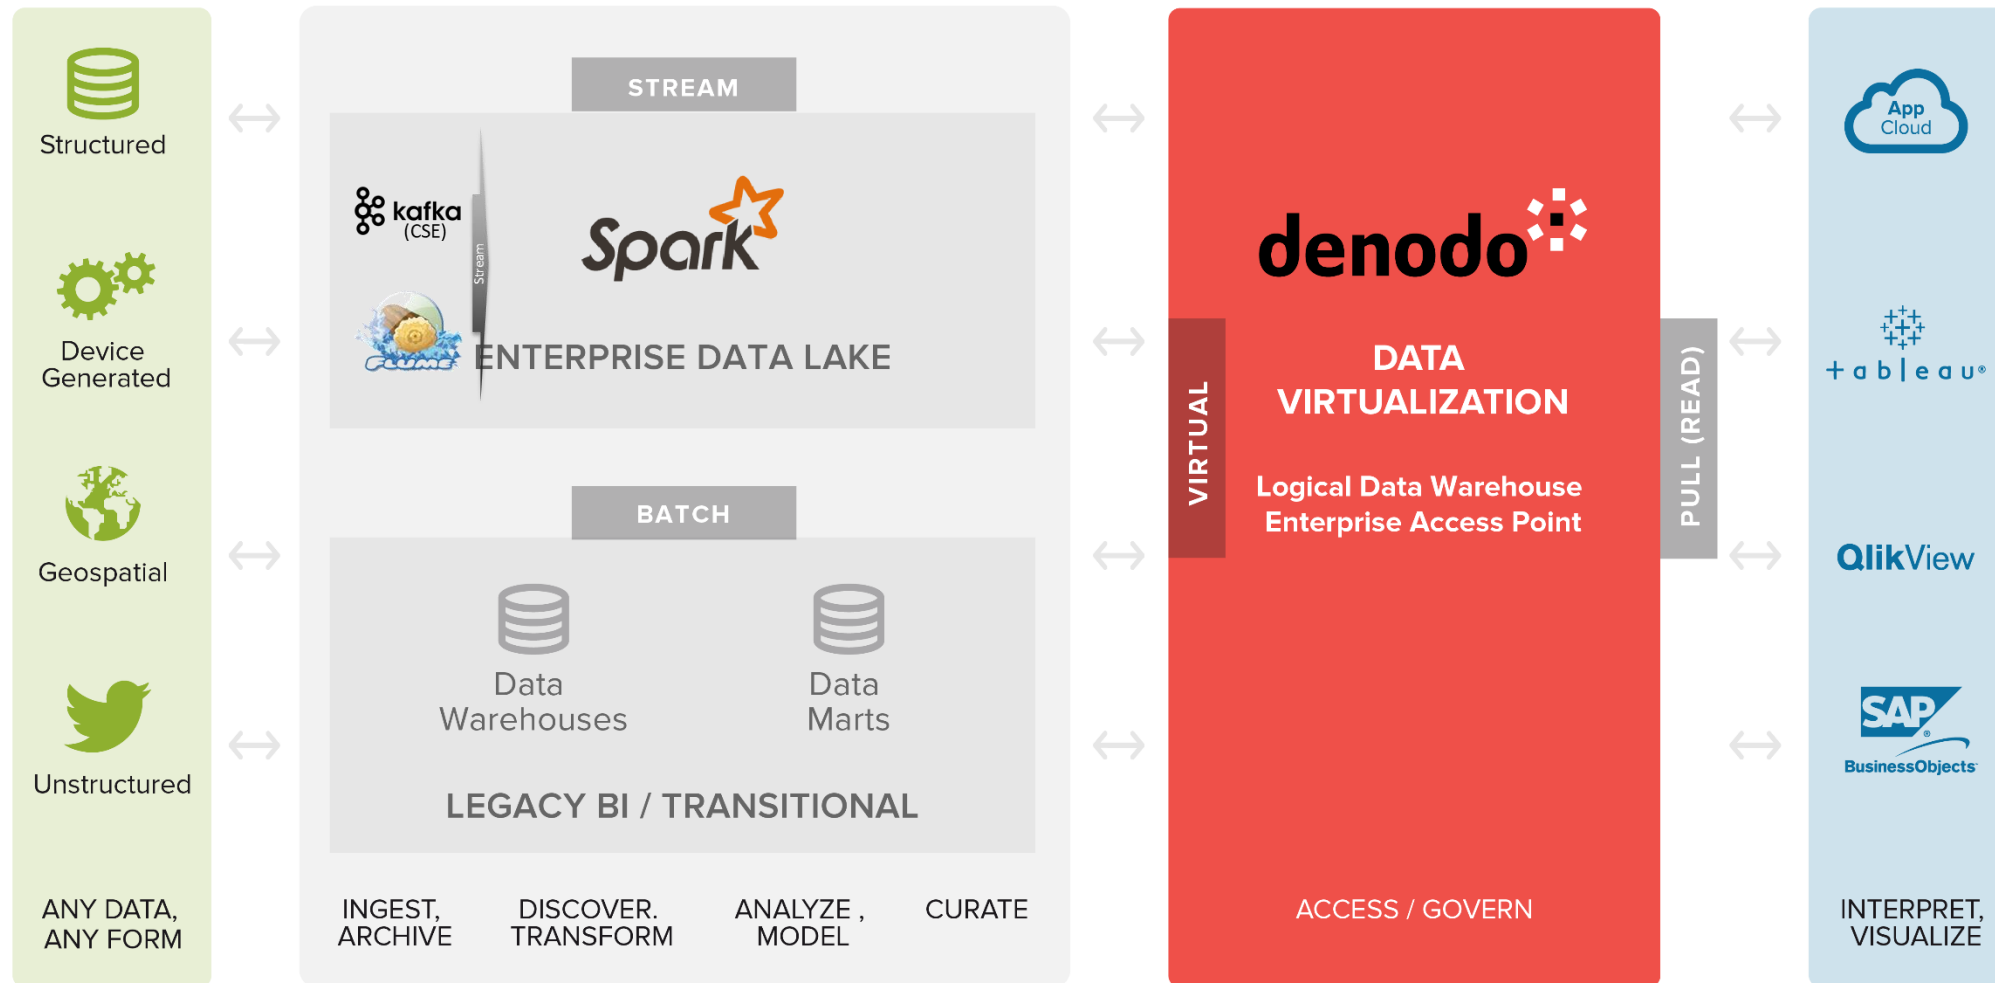
Cloud Solutions

- ✓ Cloud Modernization
- ✓ Cloud Analytics
- ✓ Hybrid Data Fabric

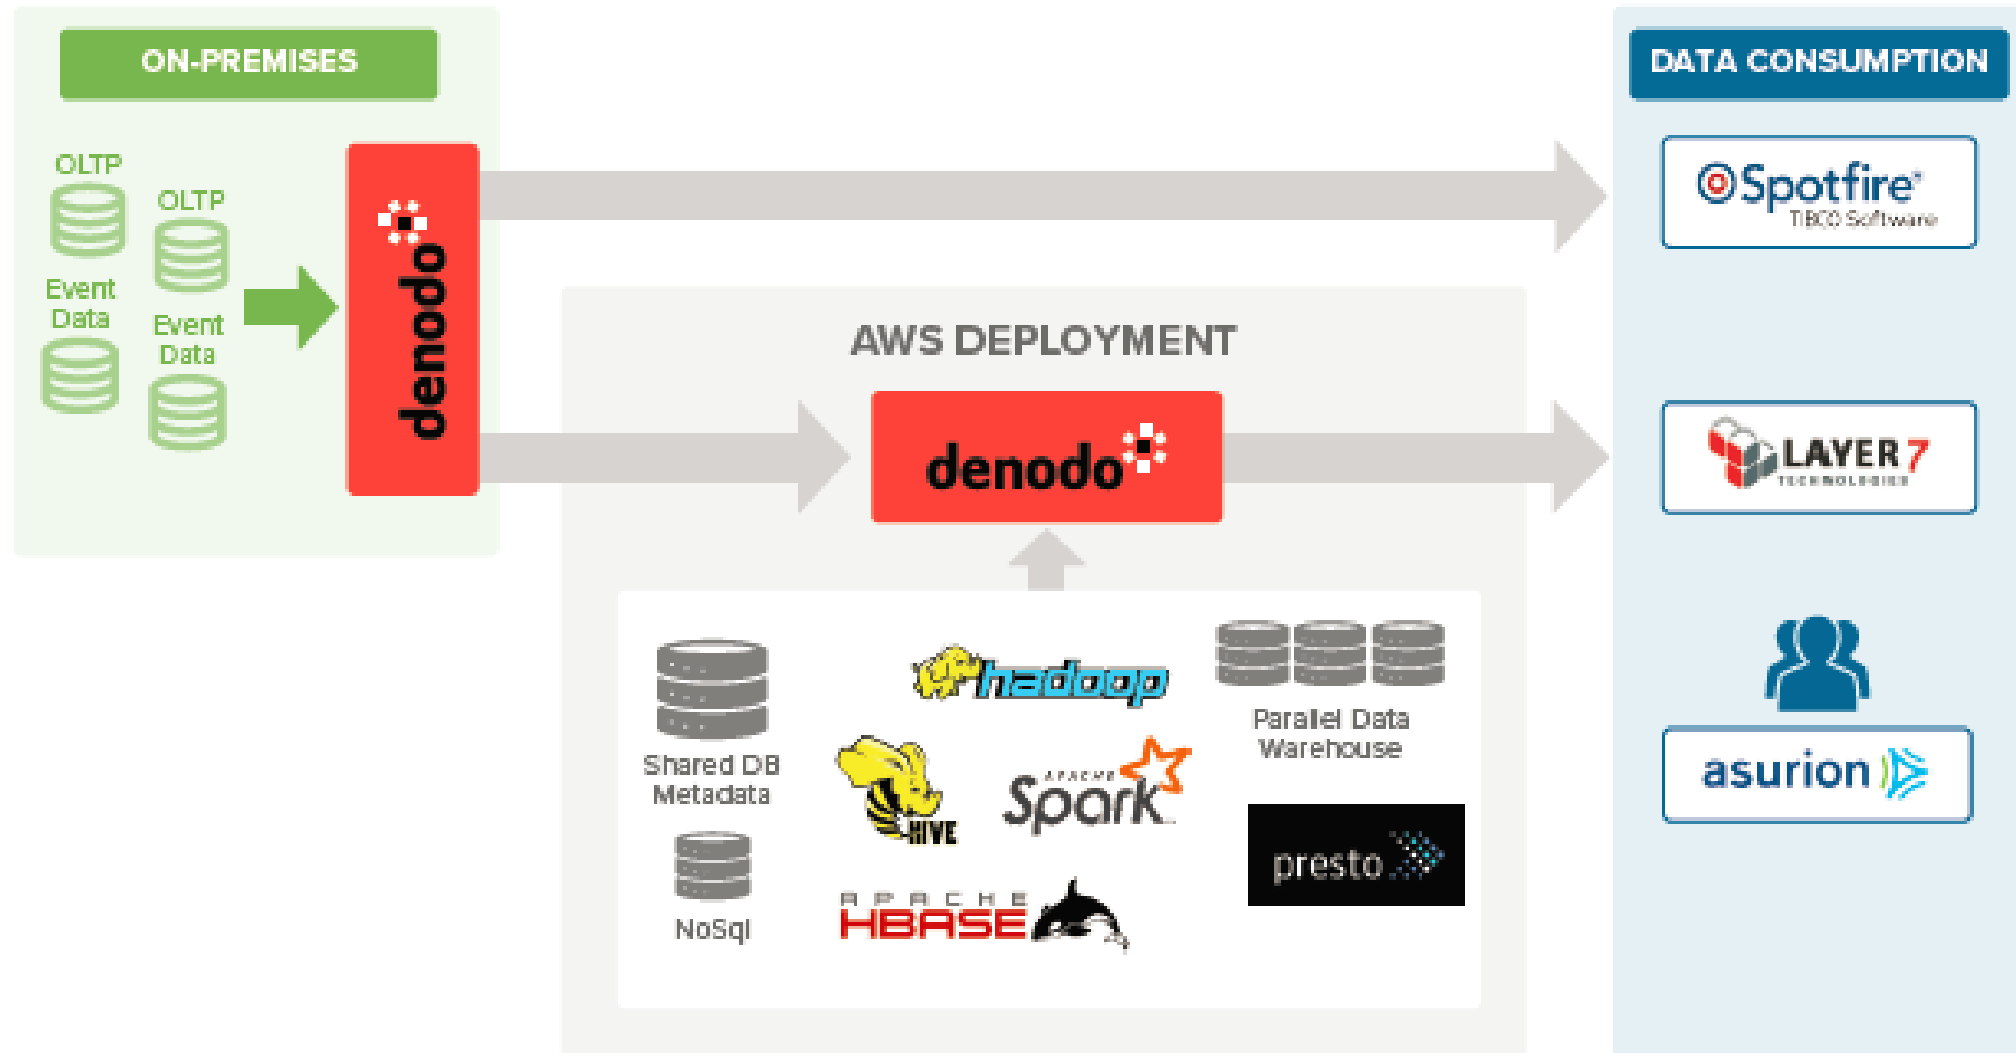
Autodesk – Integration of Streaming Data

CIO 100
AWARDS
AUTODESK
2017 CIO 100
Award

Gartner
AUTODESK
FINALIST in 2017
Excellence Awards



Asurion: Data Virtualization On-Premises and Cloud



Denodo

The Leader in Data Virtualization



LEADERSHIP

- Longest continuous focus on data virtualization – since 1999
- Leader in 2018 Forrester Wave – Big Data Fabric
- Winner of numerous awards



DENODO OFFICES, CUSTOMERS, PARTNERS

Palo Alto, CA.


Global presence throughout North America, EMEA, APAC, and Latin America.

CUSTOMERS

~500 customers, including many F500 and G2000 companies across every major industry have gained significant business agility and ROI.

FINANCIALS

Backed by \$4B+ private equity firm.
50+% annual growth; Profitable.



www.denodo.com

info@denodo.com

© Copyright Denodo Technologies. All rights reserved

Unless otherwise specified, no part of this PDF file may be reproduced or utilized in any for or by any means, electronic or mechanical, including photocopying and microfilm, without prior the written authorization from Denodo Technologies.